

SBGM: Score-Based Generative Models in JAX.

Jed Homer  1,2*

1 University Observatory, Faculty for Physics, Ludwig-Maximilians-Universität München, Scheinerstrasse
2, München, Deutschland.  2 Munich Center for Machine Learning.  * These authors contributed
equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Open Journals](#) 

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright
and release the work under a
Creative Commons Attribution 4.0
International License ([CC BY 4.0](#)).

In partnership with



This article and software are linked
with research article DOI
[10.3847/xxxxx](https://doi.org/10.3847/xxxxx) <- update this
with the DOI from AAS once you
know it., published in the
Astrophysical Journal <- The
name of the AAS journal..

Summary

Diffusion models (Ho et al., 2020; Sohl-Dickstein et al., 2015; Song, Sohl-Dickstein, et al., 2021) have emerged as the dominant paradigm for generative modelling based on performance in a variety of tasks (Peebles & Xie, 2023a; Rombach et al., 2022). The advantages of accurate density estimation and high-quality samples of normalising flows (Grathwohl et al., 2018; Papamakarios et al., 2021), VAEs (Diederik P. Kingma & Welling, 2022) and GANs (Goodfellow et al., 2014) are subsumed into this method. Significant limitations exist on implicit and neural network based likelihood models with respect to modelling normalised probability distributions and sampling speed. Score-matching diffusion models are more efficient than previous generative model algorithms for these tasks. The diffusion process is agnostic to the data representation meaning different types of data such as audio, point-clouds, videos and images can be modelled.

Statement of need

Diffusion-based generative models (Ho et al., 2020; Sohl-Dickstein et al., 2015) are a method for sampling from high-dimensional distributions. A sub-class of these models, score-based diffusion generatives models (SBGMS, (Song, Sohl-Dickstein, et al., 2021)), permit exact-likelihood estimation via a change-of-variables associated with the forward diffusion process (Song, Durkan, et al., 2021). Diffusion models allow fitting generative models to high-dimensional data in a more efficient way than normalising flows since only one neural network model parameterises the diffusion process as opposed to a sequence of neural networks in typical normalising flow architectures. Whilst existing diffusion models (Ho et al., 2020; Diederik P. Kingma et al., 2023) allow for sampling, they are limited to inaccurate variational inference approaches for density estimation which limits their use for Bayesian inference. This code provides density estimation with diffusion models using GPU enabled ODE solvers in jax (Bradbury et al., 2018) and diffraex (Kidger, 2022). Similar codes (e.g. (Rozet, 2024)) exist for diffusion models but they do not implement log-likelihood calculations, various network architectures and parallelised ODE-sampling.

The software we present, sbgm, is designed to be used by researchers in machine learning and the natural sciences for fitting diffusion models with custom architectures for their research. These models can be fit easily with multi-accelerator training and inference within the code. Typical use cases for these kinds of generative models are emulator approaches (Spurio Mancini et al., 2022), simulation-based inference (Cranmer et al., 2020), field-level inference (Andrews et al., 2023) and general inverse problems (Feng et al., 2023; Feng & Bouman, 2024; Remy et al., 2023; Song et al., 2022) (e.g. image inpainting (Song, Sohl-Dickstein, et al., 2021) and denoising (Chung et al., 2022; Daras et al., 2024)). This code allows for seamless integration of diffusion models to these applications by providing data-generating models with easy conditioning of the data on any modality. Furthermore, the implementation in equinox

⁴³ (Kidger & Garcia, 2021) guarantees safe integration of sbgm with any other sampling libraries
⁴⁴ (e.g. BlackJAX Cabezas et al. (2024)) or jax (Bradbury et al., 2018) based codes.

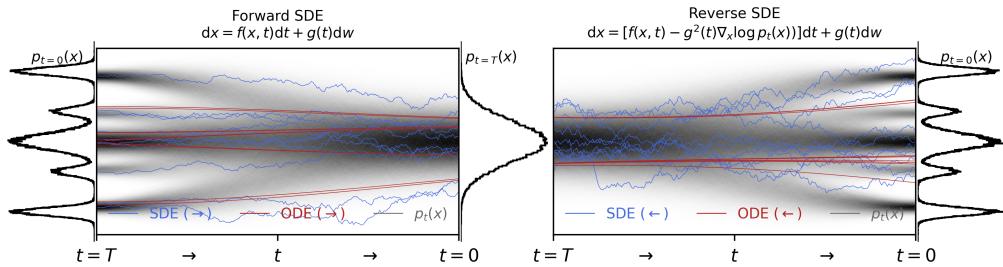


Figure 1: A diagram showing how to map data to a noise distribution (the prior) with an SDE, and reverse this SDE for generative modeling. One can also reverse the associated probability flow ODE, which yields a deterministic reverse process. Both the reverse-time SDE and probability flow ODE can be obtained by estimating the score.

⁴⁵ Diffusion

⁴⁶ Diffusion in the context of generative modelling describes the process of adding small amounts
⁴⁷ of noise sequentially to samples of data x (Sohl-Dickstein et al., 2015). A generative model
⁴⁸ for the data arises from training a neural network to reverse this process by subtracting the
⁴⁹ noise added to the data.

⁵⁰ Score-based diffusion models (Song, Sohl-Dickstein, et al., 2021) model a forward diffusion
⁵¹ process with Stochastic Differential Equations (SDEs) of the form

$$dx = f(x, t)dt + g(t)dw,$$

⁵² where $f(x, t)$ is a vector-valued function called the drift coefficient, $g(t)$ is the diffusion
⁵³ coefficient and dw is a sample of noise $dw \sim \mathcal{G}[dw|0, I]$. This equation describes the infinitely
⁵⁴ many samples of noise along the diffusion time t that perturb the data. The diffusion path,
⁵⁵ defined by the SDE, begins at $t = 0$ and ends at $T = 0$ where the resulting distribution is
⁵⁶ then a multivariate Gaussian with mean zero and covariance I .

⁵⁷ The reverse of the SDE, mapping from multivariate Gaussian samples $x(T)$ to samples of data
⁵⁸ $x(0)$, is of the form

$$dx = [f(x, t) - g^2(t)\nabla_x \log p_t(x)]dt + g(t)dw,$$

⁵⁹ where the score function $\nabla_x \log p_t(x)$ is substituted with a neural network $s_\theta(x(t), t)$ for the
⁶⁰ sampling process. The network is fit by score-matching (Hyvärinen, 2005; Vincent, 2011)
⁶¹ across the time span $[0, T]$. This network predicts the noise added to the image at time
⁶² t with the forward diffusion process, in accordance with the SDE, and removes it. With a
⁶³ data-dimensional sample of Gaussian noise from the prior $p_T(x)$ (see Figure 1) one can reverse
⁶⁴ the diffusion process to generate data.

⁶⁵ The reverse SDE may be solved with Euler-Murayama sampling (Song, Sohl-Dickstein, et al.,
⁶⁶ 2021) (or other annealed Langevin sampling methods) which is featured in the code.

⁶⁷ Likelihood calculations with diffusion models

⁶⁸ Many of the applications of generative models depend on being able to calculate the likelihood
⁶⁹ of data. Song, Sohl-Dickstein, et al. (2021) show that any SDE may be converted into an

⁷⁰ ordinary differential equation (ODE) without changing the distributions, defined by the SDE,
⁷¹ from which the noise is sampled from in the diffusion process (denoted $p_t(x)$ and shown in
⁷² grey in Figure 1). This ODE is known as the probability flow ODE (Song, Sohl-Dickstein, et
⁷³ al., 2021; Song, Durkan, et al., 2021) and is written

$$dx = [f(x, t) - g^2(t) \nabla_x \log p_t(x)] dt = f'(x, t) dt.$$

⁷⁴ This ODE can be solved with an initial-value problem. Starting with a data point $x(0) \sim p(x)$,
⁷⁵ this point is mapped along the probability flow ODE path (see the right-hand side of Figure 1)
⁷⁶ to a sample from the multivariate Gaussian prior. This inherits the formalism of continuous
⁷⁷ normalising flows (Chen et al., 2019; Grathwohl et al., 2018) without the expensive ODE
⁷⁸ simulations used to train these models - allowing for a likelihood estimate based on diffusion
⁷⁹ models (Song, Durkan, et al., 2021). The initial value problem provides a solution $x(T)$ and
⁸⁰ the change in probability along the path $\Delta = \log p(x(0)) - \log p(x(T))$ where $p(x(T))$ is a
⁸¹ simple multivariate Gaussian distribution.

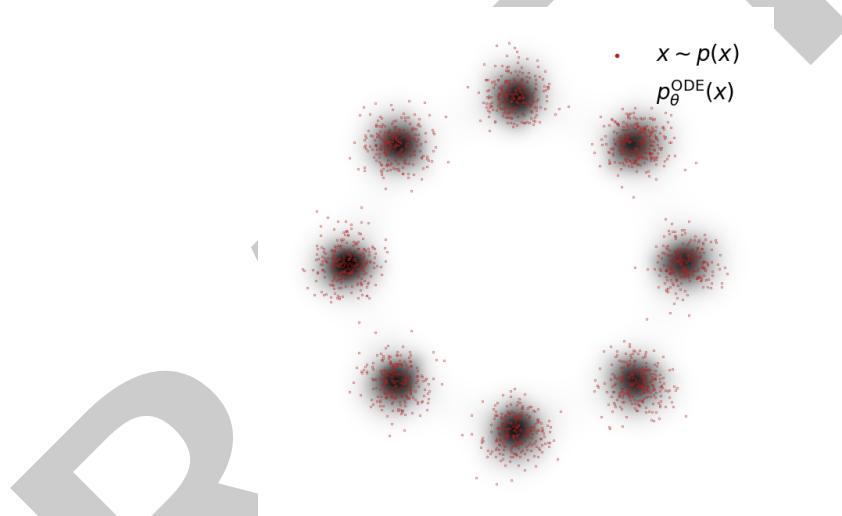


Figure 2: A diagram showing a log-likelihood calculation over the support of a Gaussian mixture model with eight components. Data is drawn (shown in red) from this mixture to train the diffusion model that gives the likelihood in gray. The log-likelihood is calculated using the ODE and a trained diffusion model.

⁸² The likelihood estimate under a score-based diffusion model is estimated by solving the
⁸³ change-of-variables equation for continuous normalising flows.

$$\frac{\partial}{\partial t} \log p(x(t)) = \nabla_x \cdot f(x(t), t),$$

⁸⁴ which gives the log-likelihood of a single datapoint $x(0)$ as

$$\log p(x(0)) = \log p(x(T)) + \int_{t=0}^{t=T} dt \nabla_x \cdot f(x, t).$$

⁸⁵ The code implements these calculations also for the Hutchinson trace estimation method and
⁸⁶ (1990) that reduces the computational expense of the estimate. Figure 2 shows an example of
⁸⁷ a data-likelihood calculation using a trained diffusion model with the ODE associated from an
⁸⁸ SDE.

89 Implementations and future work

90 Diffusion models are defined in sbgm via a score-network model s_θ and an SDE. All the
91 available SDEs (variance exploding (VE), variance preserving (VP) and sub-variance preserving
92 (SubVP) (Song, Sohl-Dickstein, et al., 2021)) in the literature of score-based diffusion models
93 are available. We provide implementations for UNet (Ronneberger et al., 2015), Diffusion
94 Transformers (Peebles & Xie, 2023b), MLP-Mixer (Tolstikhin et al., 2021) and Residual
95 Network (He et al., 2015) models which are state-of-the-art for diffusion tasks. It is possible to
96 fit score-based diffusion models to a conditional distribution $p(x|\pi, y)$ where in typical inverse
97 problems y would be an image and π a set of parameters in a physical model for the data
98 (Batzolis et al., 2021) (e.g. to solve inverse problems). The code is compatible with any model
99 written in the equinox (Kidger & Garcia, 2021) framework. We are extending the code to
100 provide transformer-based (Peebles & Xie, 2023a) and latent diffusion models (Rombach et
101 al., 2022).

102 Acknowledgements

103 We thank the developers of the packages jax (Bradbury et al., 2018), optax (DeepMind et
104 al., 2020), equinox (Kidger & Garcia, 2021) and diffrax (Kidger, 2022) for their work and
105 for making their code available to the community.

106 References

- 107 and, M. F. H. (1990). A stochastic estimator of the trace of the influence matrix for laplacian
108 smoothing splines. *Communications in Statistics - Simulation and Computation*, 19(2),
109 433–450. <https://doi.org/10.1080/03610919008812866>
- 110 Andrews, A., Jasche, J., Lavaux, G., & Schmidt, F. (2023). Bayesian field-level inference of
111 primordial non-gaussianity using next-generation galaxy surveys. *Monthly Notices of the
112 Royal Astronomical Society*, 520(4), 5746–5763. <https://doi.org/10.1093/mnras/stad432>
- 113 Batzolis, G., Stanczuk, J., Schönlieb, C.-B., & Etmann, C. (2021). *Conditional image
114 generation with score-based diffusion models*. <https://arxiv.org/abs/2111.13606>
- 115 Bradbury, J., Frostig, R., Hawkins, P., Johnson, M. J., Leary, C., Maclaurin, D., Necula, G.,
116 Paszke, A., VanderPlas, J., Wanderman-Milne, S., & Zhang, Q. (2018). JAX: Composable
117 transformations of Python+NumPy programs (Version 0.3.13). <http://github.com/jax-ml/jax>
- 118 Cabezas, A., Corenflos, A., Lao, J., & Louf, R. (2024). BlackJAX: Composable Bayesian
119 inference in JAX. <https://arxiv.org/abs/2402.10797>
- 120 Chen, R. T. Q., Rubanova, Y., Bettencourt, J., & Duvenaud, D. (2019). Neural ordinary
121 differential equations. <https://arxiv.org/abs/1806.07366>
- 122 Chung, H., Kim, J., Kim, S., & Ye, J. C. (2022). Parallel diffusion models of operator and
123 image for blind inverse problems. <https://arxiv.org/abs/2211.10656>
- 124 Cranmer, K., Brehmer, J., & Louppe, G. (2020). The frontier of simulation-based inference.
125 *Proceedings of the National Academy of Sciences*, 117(48), 30055–30062. <https://doi.org/10.1073/pnas.1912789117>
- 126 Daras, G., Dimakis, A. G., & Daskalakis, C. (2024). Consistent diffusion meets tweedie:
127 Training exact ambient diffusion models with noisy data. <https://arxiv.org/abs/2404.10177>
- 128 DeepMind, Babuschkin, I., Baumli, K., Bell, A., Bhupatiraju, S., Bruce, J., Buchlovsky, P.,
129 Budden, D., Cai, T., Clark, A., Danihelka, I., Dedieu, A., Fantacci, C., Godwin, J., Jones,
130 Homer. (2025). SBGM: Score-Based Generative Models in JAX. *Journal of Open Source Software*, VOL? (ISSUE?), PAGE? <https://doi.org/10.4xxxxx/draft>

- 132 C., Hemsley, R., Hennigan, T., Hessel, M., Hou, S., ... Viola, F. (2020). *The DeepMind*
133 *JAX Ecosystem*. <http://github.com/google-deepmind>
- 134 Feng, B. T., & Bouman, K. L. (2024). *Variational bayesian imaging with an efficient surrogate*
135 *score-based prior*. <https://arxiv.org/abs/2309.01949>
- 136 Feng, B. T., Smith, J., Rubinstein, M., Chang, H., Bouman, K. L., & Freeman, W. T.
137 (2023). *Score-based diffusion models as principled priors for inverse imaging*. <https://arxiv.org/abs/2304.11751>
- 138 Goodfellow, I. J., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., Courville,
139 A., & Bengio, Y. (2014). *Generative adversarial networks*. <https://arxiv.org/abs/1406.2661>
- 140 Grathwohl, W., Chen, R. T. Q., Bettencourt, J., Sutskever, I., & Duvenaud, D. (2018).
141 *FFJORD: Free-form continuous dynamics for scalable reversible generative models*. <https://arxiv.org/abs/1810.01367>
- 142 He, K., Zhang, X., Ren, S., & Sun, J. (2015). *Deep residual learning for image recognition*.
143 <https://arxiv.org/abs/1512.03385>
- 144 Ho, J., Jain, A., & Abbeel, P. (2020). *Denoising diffusion probabilistic models*. <https://arxiv.org/abs/2006.11239>
- 145 Hyvärinen, A. (2005). Estimation of non-normalized statistical models by score matching.
146 *Journal of Machine Learning Research*, 6(24), 695–709. <http://jmlr.org/papers/v6/hyvarinen05a.html>
- 147 Kidger, P. (2022). *On neural differential equations*. <https://arxiv.org/abs/2202.02435>
- 148 Kidger, P., & Garcia, C. (2021). Equinox: Neural networks in JAX via callable PyTrees and
149 filtered transformations. *Differentiable Programming Workshop at Neural Information
150 Processing Systems 2021*.
- 151 Kingma, Diederik P., Salimans, T., Poole, B., & Ho, J. (2023). *Variational diffusion models*.
152 <https://arxiv.org/abs/2107.00630>
- 153 Kingma, Diederik P., & Welling, M. (2022). *Auto-encoding variational bayes*. <https://arxiv.org/abs/1312.6114>
- 154 Papamakarios, G., Nalisnick, E., Rezende, D. J., Mohamed, S., & Lakshminarayanan, B.
155 (2021). *Normalizing flows for probabilistic modeling and inference*. <https://arxiv.org/abs/1912.02762>
- 156 Peebles, W., & Xie, S. (2023a). *Scalable diffusion models with transformers*. <https://arxiv.org/abs/2212.09748>
- 157 Peebles, W., & Xie, S. (2023b). *Scalable diffusion models with transformers*. <https://arxiv.org/abs/2212.09748>
- 158 Remy, B., Lanusse, F., Jeffrey, N., Liu, J., Starck, J.-L., Osato, K., & Schrabbach, T. (2023).
159 Probabilistic mass-mapping with neural score estimation. *Astronomy & Astrophysics*,
160 672, A51. <https://doi.org/10.1051/0004-6361/202243054>
- 161 Rombach, R., Blattmann, A., Lorenz, D., Esser, P., & Ommer, B. (2022). *High-resolution*
162 *image synthesis with latent diffusion models*. <https://arxiv.org/abs/2112.10752>
- 163 Ronneberger, O., Fischer, P., & Brox, T. (2015). *U-net: Convolutional networks for biomedical*
164 *image segmentation*. <https://arxiv.org/abs/1505.04597>
- 165 Rozet, F. (2024). *Azula: Diffusion models in PyTorch* (Version 0.2.04). <https://github.com/probabilists/azula>
- 166 Sohl-Dickstein, J., Weiss, E. A., Maheswaranathan, N., & Ganguli, S. (2015). *Deep unsupervised*
167 *learning using nonequilibrium thermodynamics*. <https://arxiv.org/abs/1503.03585>
- 168

- 177 Song, Y., Durkan, C., Murray, I., & Ermon, S. (2021). *Maximum likelihood training of*
178 *score-based diffusion models*. <https://arxiv.org/abs/2101.09258>
- 179 Song, Y., Shen, L., Xing, L., & Ermon, S. (2022). *Solving inverse problems in medical imaging*
180 *with score-based generative models*. <https://arxiv.org/abs/2111.08005>
- 181 Song, Y., Sohl-Dickstein, J., Kingma, D. P., Kumar, A., Ermon, S., & Poole, B. (2021).
182 *Score-based generative modeling through stochastic differential equations*. <https://arxiv.org/abs/2011.13456>
- 184 Spurio Mancini, A., Piras, D., Alsing, J., Joachimi, B., & Hobson, M. P. (2022). <Scp>Cos-
185 *moPower</scp>*: Emulating cosmological power spectra for accelerated bayesian inference
186 from next-generation surveys. *Monthly Notices of the Royal Astronomical Society*, 511(2),
187 1771–1788. <https://doi.org/10.1093/mnras/stac064>
- 188 Tolstikhin, I., Houlsby, N., Kolesnikov, A., Beyer, L., Zhai, X., Unterthiner, T., Yung, J.,
189 Steiner, A., Keysers, D., Uszkoreit, J., Lucic, M., & Dosovitskiy, A. (2021). *MLP-mixer:*
190 *An all-MLP architecture for vision*. <https://arxiv.org/abs/2105.01601>
- 191 Vincent, P. (2011). A connection between score matching and denoising autoencoders. *Neural*
192 *Computation*, 23(7), 1661–1674. https://doi.org/10.1162/NECO_a_00142

DRAFT