

SBGM: Score-Based Generative Models in JAX.

Jed Homer^{1,2*}

¹ University Observatory, Faculty for Physics, Ludwig-Maximilians-Universität München, Scheinerstrasse 1, München, Deutschland. ² Munich Center for Machine Learning. * These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [Open Journals](#)

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

In partnership with



This article and software are linked with research article DOI [10.3847/xxxx <- update this with the DOI from AAS once you know it.](https://doi.org/10.3847/xxxx <- update this with the DOI from AAS once you know it.), published in the Astrophysical Journal <- The name of the AAS journal..

Summary

Diffusion models ([Ho et al., 2020](#); [Sohl-Dickstein et al., 2015](#); [Song, Sohl-Dickstein, et al., 2021](#)) have emerged as the dominant paradigm for generative modeling based on performance in a variety of tasks ([Peebles & Xie, 2023](#); [Rombach et al., 2022](#)). The advantages of accurate density estimation and high-quality samples of normalising flows ([Grathwohl et al., 2018](#); [Papamakarios et al., 2021](#)), VAEs ([Diederik P. Kingma & Welling, 2022](#)) and GANs ([Goodfellow et al., 2014](#)) are subsumed into this method. Significant limitations exist on implicit and neural network based likelihood models with respect to modeling normalised probability distributions and sampling speed. Score-matching diffusion models are more efficient than previous generative model algorithms for these tasks. The diffusion process is agnostic to the data representation meaning different types of data such as audio, point-clouds, videos and images can be modelled.

Statement of need

Diffusion-based generative models ([Ho et al., 2020](#); [Sohl-Dickstein et al., 2015](#)) are a method for sampling from high-dimensional distributions. A sub-class of these models, score-based diffusion generative models (SBGMs, ([Song, Sohl-Dickstein, et al., 2021](#))), permit exact-likelihood estimation via a change-of-variables associated with the forward diffusion process ([Song, Durkan, et al., 2021](#)). Diffusion models allow fitting generative models to high-dimensional data in a more efficient way than normalising flows, since only one neural network model parameterises the diffusion process, as opposed to a sequence of neural networks in typical normalising flow architectures. Whilst existing diffusion models ([Ho et al., 2020](#); [Diederik P. Kingma et al., 2023](#)) allow for sampling, they are limited to inaccurate variational inference approaches for density estimation which limits their use for Bayesian inference. This code provides density estimation with diffusion models using GPU enabled ODE solvers in jax ([Bradbury et al., 2018](#)) and diffrafx ([Kidger, 2022](#)). Similar codes (e.g. ([Rozet, 2024](#))) exist for diffusion models but they do not implement log-likelihood calculations, various network architectures, and parallelised computations for optimisation and SDE/ODE-sampling.

The software we present, sbgm, is designed to be used by researchers in machine learning and the natural sciences for fitting diffusion models with custom architectures for their research. These models can be fit easily with multi-accelerator training and inference routines within the code (with demonstration examples provided). Typical use cases for these kinds of generative models are emulator approaches ([Spurio Mancini et al., 2022](#)), simulation-based inference ([Cranmer et al., 2020](#)), field-level inference ([Andrews et al., 2023](#)) and general inverse problems ([Feng et al., 2023](#); [Feng & Bouman, 2024](#); [Remy et al., 2023](#); [Song et al., 2022](#)) (e.g. image inpainting ([Song, Sohl-Dickstein, et al., 2021](#)) and denoising ([Chung et al., 2022](#); [Daras et al., 2024](#))). This code allows for seamless integration of diffusion models to these applications by providing data-generating models with easy conditioning of the data on any modality (e.g. images, audio

⁴³ or model parameters). Furthermore, the implementation in equinox ([Kidger & Garcia, 2021](#))
⁴⁴ guarantees safe integration of sbgm with any other sampling libraries (e.g. BlackJAX ([Cabezas et al., 2024](#))) or jax ([Bradbury et al., 2018](#)) based codes.

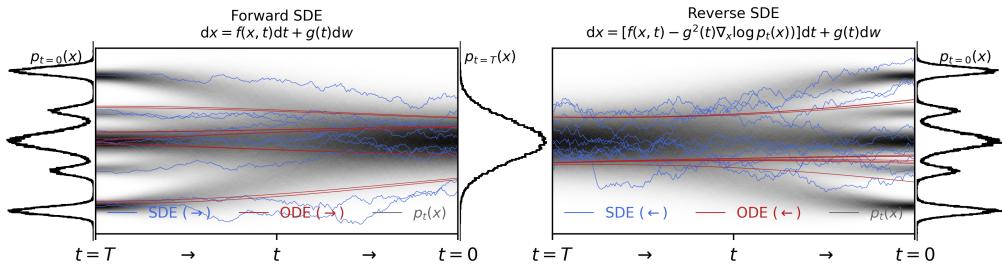


Figure 1: A diagram showing how to map data to a noise distribution (the prior) with an SDE, and reverse this SDE for generative modeling. One can also reverse the associated probability flow ODE, which yields a deterministic reverse process. Both the reverse-time SDE and probability flow ODE can be obtained by estimating the score.

⁴⁶ Diffusion

⁴⁷ Diffusion in the context of generative modeling describes the process of adding small amounts
⁴⁸ of noise sequentially to samples of data x ([Sohl-Dickstein et al., 2015](#)). A generative model
⁴⁹ for the data arises from training a neural network to reverse this process by subtracting the
⁵⁰ noise added to the data.

⁵¹ Score-based diffusion models ([Song, Sohl-Dickstein, et al., 2021](#)) model a forward diffusion
⁵² process with Stochastic Differential Equations (SDEs) of the form

$$dx_t = f(x_t, t)dt + g(t)d\omega_t,$$

⁵³ where $f(x_t, t)$ is a vector-valued function called the drift coefficient, $g(t)$ is the diffusion
⁵⁴ coefficient and $d\omega_t$ is a sample of noise $d\omega_t \sim \mathcal{G}[d\omega_t | \mathbf{0}, \mathbf{I}]$. This equation describes the
⁵⁵ infinitely many samples of noise along the diffusion time t that perturb the data. The diffusion
⁵⁶ path, defined by the SDE, begins at $t = 0$ and ends at $t = T$ where the resulting distribution is
⁵⁷ then a multivariate Gaussian with mean zero and covariance \mathbf{I} . The code implements various
⁵⁸ SDEs known in the diffusion model literature.

⁵⁹ The reverse of the SDE, mapping from multivariate Gaussian samples x_T to samples of data
⁶⁰ x_0 , is of the form

$$dx_t = [f(x_t, t) - g^2(t)\nabla_{x_t} \log p_t(x_t)]dt + g(t)d\omega_t,$$

⁶¹ where the score function $\nabla_{x_t} \log p_t(x_t)$ is substituted with a neural network $s_\theta(x(t), t)$ for
⁶² the sampling process. The network is fit by score-matching ([Hyvärinen, 2005; Vincent, 2011](#))
⁶³ across the time span $[0, T]$. This network predicts the noise added to the image at time
⁶⁴ t with the forward diffusion process, in accordance with the SDE, and removes it. With a
⁶⁵ data-dimensional sample of Gaussian noise from the prior $p_T(x_T)$ (see Figure 1) one can
⁶⁶ reverse the diffusion process to generate data.

⁶⁷ The reverse SDE may be solved with Euler-Maruyama sampling ([Song, Sohl-Dickstein, et al., 2021](#))
⁶⁸ (or other annealed Langevin sampling methods) which is featured in the code.

69 Likelihood calculations with diffusion models

70 Many of the applications of generative models depend on being able to calculate the likelihood
 71 of data. Song, Sohl-Dickstein, et al. (2021) show that any SDE may be converted into an
 72 ordinary differential equation (ODE) without changing the distributions $p_t(\mathbf{x}_t)$, defined by the
 73 SDE, from which the noise is sampled from in the diffusion process (denoted $p_t(\mathbf{x}_t)$ and shown
 74 in grey in Figure 1). This ODE is known as the probability flow ODE (Song, Sohl-Dickstein,
 75 et al., 2021; Song, Durkan, et al., 2021) and is written

$$d\mathbf{x}_t = [f(\mathbf{x}_t, t) - g^2(t)\nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t)]dt = f'(\mathbf{x}_t, t)dt.$$

76 This ODE can be solved with an initial-value problem to sample new data or estimate its
 77 density. Starting with a data point $\mathbf{x}_0 \sim p(\mathbf{x}) = p_0(\mathbf{x}_0)$, this point is mapped along the
 78 probability flow ODE path (see the right-hand side of Figure 1) shows a mapping to a sample
 79 from the multivariate Gaussian prior $\mathbf{x}_T \sim p_T(\mathbf{x}_T)$. This inherits the formalism of continuous
 80 normalising flows (Chen et al., 2019; Grathwohl et al., 2018) without the expensive ODE
 81 simulations used to train these models - allowing for a likelihood estimate based on diffusion
 82 models (Song, Durkan, et al., 2021). The initial value problem provides a solution \mathbf{x}_T and the
 83 change in probability along the path $\Delta = \log p_0(\mathbf{x}_0) - \log p_T(\mathbf{x}_T)$ where $p_T(\mathbf{x}_T)$ is a simple
 84 multivariate Gaussian distribution. Various ODE solvers of different orders are available (for
 85 a user to balance speed and accuracy of sampling) which are provided by diffraax (Kidger,
 86 2022).

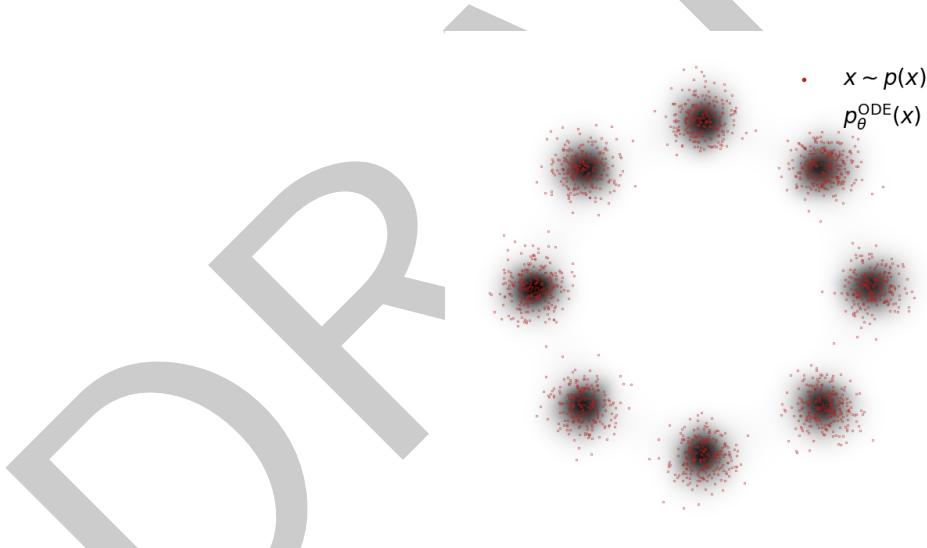


Figure 2: A diagram showing a log-likelihood calculation over the support of a Gaussian mixture model with eight components. Data is drawn (shown in red) from this mixture to train the diffusion model that gives the likelihood (defined by the diffusion model) in gray. The log-likelihood is calculated using the ODE and a trained diffusion model.

87 The likelihood estimate under a score-based diffusion model is estimated by solving the
 88 change-of-variables equation for continuous normalising flows

$$\frac{\partial}{\partial t} \log p_t(\mathbf{x}_t) = -\nabla_{\mathbf{x}_t} \cdot f(\mathbf{x}_t, t),$$

89 which gives the log-likelihood of a single datapoint \mathbf{x}_0 as

$$\log p(\mathbf{x}_0) = \log p(\mathbf{x}_T) - \int_{t=0}^{t=T} dt \nabla_{\mathbf{x}_t} \cdot f(\mathbf{x}_t, t).$$

90 The code implements these calculations also for the Hutchinson trace estimation method
 91 ([Grathwohl et al., 2018](#); [Hutchinson, 1990](#)) which reduces the computational expense of the
 92 density estimate. This is because the divergence term $\nabla_{\mathbf{x}_t} \cdot f(\mathbf{x}_t, t)$ is materialised via a
 93 cheaper vector-Jacobian product (i.e. a $\mathcal{O}(\dim(\mathbf{x}))$ operation versus a $\mathcal{O}(\dim(\mathbf{x})^2)$ operation
 94 for the full Jacobian). Figure 2 shows an example of a data-likelihood calculation using a
 95 trained diffusion model, where the density estimate from the ODE is calculated after training
 96 with the associated SDE of the forward diffusion process.

97 **Implementations and future work**

98 Diffusion models are defined in sbgm via a score-network model s_θ and an SDE. All the
 99 available SDEs (variance exploding (VE), variance preserving (VP) and sub-variance preserving
 100 (SubVP) ([Song, Sohl-Dickstein, et al., 2021](#))) in the literature of score-based diffusion models
 101 are available. We provide implementations for UNet ([Ronneberger et al., 2015](#)), Diffusion
 102 Transformers ([Peebles & Xie, 2023](#)), MLP-Mixer ([Tolstikhin et al., 2021](#)) and Residual Network
 103 ([He et al., 2015](#)) models which are state-of-the-art for diffusion tasks. It is possible to fit
 104 score-based diffusion models to a conditional distribution $p(\mathbf{x}|\pi, \mathbf{y})$ where in typical inverse
 105 problems \mathbf{y} would be an image and π a set of parameters in a physical model for the data
 106 ([Batzolis et al., 2021](#)) (e.g. to solve inverse problems). The code is compatible with any model
 107 written in the equinox ([Kidger & Garcia, 2021](#)) framework. We recently extended the code to
 108 provide transformer-based diffusion models ([Peebles & Xie, 2023](#)) and plan to extend to latent
 109 diffusion models ([Rombach et al., 2022](#)) and flow matching ([Lipman et al., 2023](#)).

110 **Acknowledgements**

111 We thank the developers of the packages jax ([Bradbury et al., 2018](#)), optax ([DeepMind et
 112 al., 2020](#)), equinox ([Kidger & Garcia, 2021](#)) and diffrafx ([Kidger, 2022](#)) for their work and
 113 for making their code available to the community.

114 **References**

- 115 Andrews, A., Jasche, J., Lavaux, G., & Schmidt, F. (2023). Bayesian field-level inference of
 116 primordial non-gaussianity using next-generation galaxy surveys. *Monthly Notices of the
 117 Royal Astronomical Society*, 520(4), 5746–5763. <https://doi.org/10.1093/mnras/stad432>
- 118 Batzolis, G., Stanczuk, J., Schönlieb, C.-B., & Etmann, C. (2021). *Conditional image
 119 generation with score-based diffusion models*. <https://arxiv.org/abs/2111.13606>
- 120 Bradbury, J., Frostig, R., Hawkins, P., Johnson, M. J., Leary, C., Maclaurin, D., Necula, G.,
 121 Paszke, A., VanderPlas, J., Wanderman-Milne, S., & Zhang, Q. (2018). *JAX: Composable
 122 transformations of Python+NumPy programs* (Version 0.3.13). <http://github.com/jax-ml/jax>
- 124 Cabezas, A., Corenflos, A., Lao, J., & Louf, R. (2024). *BlackJAX: Composable Bayesian
 125 inference in JAX*. <https://arxiv.org/abs/2402.10797>
- 126 Chen, R. T. Q., Rubanova, Y., Bettencourt, J., & Duvenaud, D. (2019). *Neural ordinary
 127 differential equations*. <https://arxiv.org/abs/1806.07366>
- 128 Chung, H., Kim, J., Kim, S., & Ye, J. C. (2022). *Parallel diffusion models of operator and
 129 image for blind inverse problems*. <https://arxiv.org/abs/2211.10656>

- 130 Cranmer, K., Brehmer, J., & Louppe, G. (2020). The frontier of simulation-based inference.
131 *Proceedings of the National Academy of Sciences*, 117(48), 30055–30062. <https://doi.org/10.1073/pnas.1912789117>
- 133 Daras, G., Dimakis, A. G., & Daskalakis, C. (2024). *Consistent diffusion meets tweedie: Training exact ambient diffusion models with noisy data.* <https://arxiv.org/abs/2404.10177>
- 135 DeepMind, Babuschkin, I., Baumli, K., Bell, A., Bhupatiraju, S., Bruce, J., Buchlovsky, P.,
136 Budden, D., Cai, T., Clark, A., Danihelka, I., Dedieu, A., Fantacci, C., Godwin, J., Jones,
137 C., Hemsley, R., Hennigan, T., Hessel, M., Hou, S., ... Viola, F. (2020). *The DeepMind
138 JAX Ecosystem.* <http://github.com/google-deepmind>
- 139 Feng, B. T., & Bouman, K. L. (2024). *Variational Bayesian Imaging with an Efficient Surrogate
140 Score-based Prior.* <https://arxiv.org/abs/2309.01949>
- 141 Feng, B. T., Smith, J., Rubinstein, M., Chang, H., Bouman, K. L., & Freeman, W. T.
142 (2023). *Score-based diffusion models as principled priors for inverse imaging.* <https://arxiv.org/abs/2304.11751>
- 144 Goodfellow, I. J., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., Courville,
145 A., & Bengio, Y. (2014). *Generative adversarial networks.* <https://arxiv.org/abs/1406.2661>
- 146 Grathwohl, W., Chen, R. T. Q., Bettencourt, J., Sutskever, I., & Duvenaud, D. (2018).
147 *FFJORD: Free-form continuous dynamics for scalable reversible generative models.* <https://arxiv.org/abs/1810.01367>
- 149 He, K., Zhang, X., Ren, S., & Sun, J. (2015). *Deep residual learning for image recognition.*
150 <https://arxiv.org/abs/1512.03385>
- 151 Ho, J., Jain, A., & Abbeel, P. (2020). *Denoising diffusion probabilistic models.* <https://arxiv.org/abs/2006.11239>
- 153 Hutchinson, M. F. (1990). A stochastic estimator of the trace of the influence matrix for
154 laplacian smoothing splines. *Communications in Statistics - Simulation and Computation*,
155 19(2), 433–450. <https://doi.org/10.1080/03610919008812866>
- 156 Hyvärinen, A. (2005). Estimation of non-normalized statistical models by score matching.
157 *Journal of Machine Learning Research*, 6(24), 695–709. <http://jmlr.org/papers/v6/hyvarinen05a.html>
- 159 Kidger, P. (2022). *On neural differential equations.* <https://arxiv.org/abs/2202.02435>
- 160 Kidger, P., & Garcia, C. (2021). Equinox: Neural networks in JAX via callable PyTrees and
161 filtered transformations. *Differentiable Programming Workshop at Neural Information
162 Processing Systems 2021.*
- 163 Kingma, Diederik P., Salimans, T., Poole, B., & Ho, J. (2023). *Variational diffusion models.*
164 <https://arxiv.org/abs/2107.00630>
- 165 Kingma, Diederik P., & Welling, M. (2022). *Auto-Encoding Variational Bayes.* <https://arxiv.org/abs/1312.6114>
- 167 Lipman, Y., Chen, R. T. Q., Ben-Hamu, H., Nickel, M., & Le, M. (2023). *Flow matching for
168 generative modeling.* <https://arxiv.org/abs/2210.02747>
- 169 Papamakarios, G., Nalisnick, E., Rezende, D. J., Mohamed, S., & Lakshminarayanan, B.
170 (2021). *Normalizing flows for probabilistic modeling and inference.* <https://arxiv.org/abs/1912.02762>
- 172 Peebles, W., & Xie, S. (2023). *Scalable diffusion models with transformers.* <https://arxiv.org/abs/2212.09748>
- 174 Remy, B., Lanusse, F., Jeffrey, N., Liu, J., Starck, J.-L., Osato, K., & Schrabbach, T. (2023).
175 Probabilistic mass-mapping with neural score estimation. *Astronomy & Astrophysics*,

- 176 672, A51. <https://doi.org/10.1051/0004-6361/202243054>
- 177 Rombach, R., Blattmann, A., Lorenz, D., Esser, P., & Ommer, B. (2022). *High-resolution*
178 *image synthesis with latent diffusion models.* <https://arxiv.org/abs/2112.10752>
- 179 Ronneberger, O., Fischer, P., & Brox, T. (2015). *U-net: Convolutional networks for biomedical*
180 *image segmentation.* <https://arxiv.org/abs/1505.04597>
- 181 Rozet, F. (2024). *Azula: Diffusion models in PyTorch* (Version 0.2.04). <https://github.com/probabilists/azula>
- 183 Sohl-Dickstein, J., Weiss, E. A., Maheswaranathan, N., & Ganguli, S. (2015). *Deep unsuper-*
184 *vised learning using nonequilibrium thermodynamics.* <https://arxiv.org/abs/1503.03585>
- 185 Song, Y., Durkan, C., Murray, I., & Ermon, S. (2021). *Maximum likelihood training of*
186 *score-based diffusion models.* <https://arxiv.org/abs/2101.09258>
- 187 Song, Y., Shen, L., Xing, L., & Ermon, S. (2022). *Solving inverse problems in medical imaging*
188 *with score-based generative models.* <https://arxiv.org/abs/2111.08005>
- 189 Song, Y., Sohl-Dickstein, J., Kingma, D. P., Kumar, A., Ermon, S., & Poole, B. (2021).
190 *Score-based generative modeling through stochastic differential equations.* <https://arxiv.org/abs/2011.13456>
- 192 Spurio Mancini, A., Piras, D., Alsing, J., Joachimi, B., & Hobson, M. P. (2022). Cos-
193 moPower: Emulating cosmological power spectra for accelerated Bayesian inference from
194 next-generation surveys. *Monthly Notices of the Royal Astronomical Society*, 511(2),
195 1771–1788. <https://doi.org/10.1093/mnras/stac064>
- 196 Tolstikhin, I., Houlsby, N., Kolesnikov, A., Beyer, L., Zhai, X., Unterthiner, T., Yung, J.,
197 Steiner, A., Keysers, D., Uszkoreit, J., Lucic, M., & Dosovitskiy, A. (2021). *MLP-mixer:*
198 *An all-MLP architecture for vision.* <https://arxiv.org/abs/2105.01601>
- 199 Vincent, P. (2011). A connection between score matching and denoising autoencoders. *Neural*
200 *Computation*, 23(7), 1661–1674. https://doi.org/10.1162/NECO_a_00142