

SBGM: Score-Based Generative Models in JAX.

Jed Homer  1,2*

1 University Observatory, Faculty for Physics, Ludwig-Maximilians-Universität München, Scheinerstrasse
2, München, Deutschland.  2 Munich Center for Machine Learning.  * These authors contributed
equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Open Journals](#) 

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright
and release the work under a
Creative Commons Attribution 4.0
International License ([CC BY 4.0](#)).

In partnership with



This article and software are linked
with research article DOI
[10.3847/xxxxx](https://doi.org/10.3847/xxxxx) <- update this
with the DOI from AAS once you
know it., published in the
Astrophysical Journal <- The
name of the AAS journal..

Summary

Diffusion models (Ho et al., 2020; Sohl-Dickstein et al., 2015; Song, Sohl-Dickstein, et al., 2021) have emerged as the dominant paradigm for generative modelling based on performance in a variety of tasks (Peebles & Xie, 2023a; Rombach et al., 2022). The advantages of accurate density estimation and high-quality samples of normalising flows (Grathwohl et al., 2018; Papamakarios et al., 2021), VAEs (Diederik P. Kingma & Welling, 2022) and GANs (Goodfellow et al., 2014) are subsumed into this method. Significant limitations exist on implicit and neural network based likelihood models with respect to modelling normalised probability distributions and sampling speed. Score-matching diffusion models are more efficient than previous generative model algorithms for these tasks. The diffusion process is agnostic to the data representation meaning different types of data such as audio, point-clouds, videos and images can be modelled.

Statement of need

Diffusion-based generative models (Ho et al., 2020; Sohl-Dickstein et al., 2015) are a method for sampling from high-dimensional distributions. A sub-class of these models, score-based diffusion generatives models (SBGMs, (Song, Sohl-Dickstein, et al., 2021)), permit exact-likelihood estimation via a change-of-variables associated with the forward diffusion process (Song, Durkan, et al., 2021). Diffusion models allow fitting generative models to high-dimensional data in a more efficient way than normalising flows since only one neural network model parameterises the diffusion process as opposed to a sequence of neural networks in typical normalising flow architectures. Whilst existing diffusion models (Ho et al., 2020; Diederik P. Kingma et al., 2023) allow for sampling, they are limited to inaccurate variational inference approaches for density estimation which limits their use for Bayesian inference. This code provides density estimation with diffusion models using GPU enabled ODE solvers in jax (Bradbury et al., 2018) and diffraex (Kidger, 2022). Similar codes (e.g. (Rozet, 2024)) exist for diffusion models but they do not implement log-likelihood calculations, various network architectures and parallelised computations for optimisation and SDE/ODE-sampling.

The software we present, sbgm, is designed to be used by researchers in machine learning and the natural sciences for fitting diffusion models with custom architectures for their research. These models can be fit easily with multi-accelerator training and inference routines within the code (with demonstration examples provided). Typical use cases for these kinds of generative models are emulator approaches (Spurio Mancini et al., 2022), simulation-based inference (Cranmer et al., 2020), field-level inference (Andrews et al., 2023) and general inverse problems (Feng et al., 2023; Feng & Bouman, 2024; Remy et al., 2023; Song et al., 2022) (e.g. image inpainting (Song, Sohl-Dickstein, et al., 2021) and denoising (Chung et al., 2022; Daras et al., 2024)). This code allows for seamless integration of diffusion models to these applications by providing data-generating models with easy conditioning of the data on any modality (e.g. images, audio

43 or model parameters). Furthermore, the implementation in equinox (Kidger & Garcia, 2021)
44 guarantees safe integration of sbgm with any other sampling libraries (e.g. BlackJAX Cabezas
45 et al. (2024)) or jax (Bradbury et al., 2018) based codes.

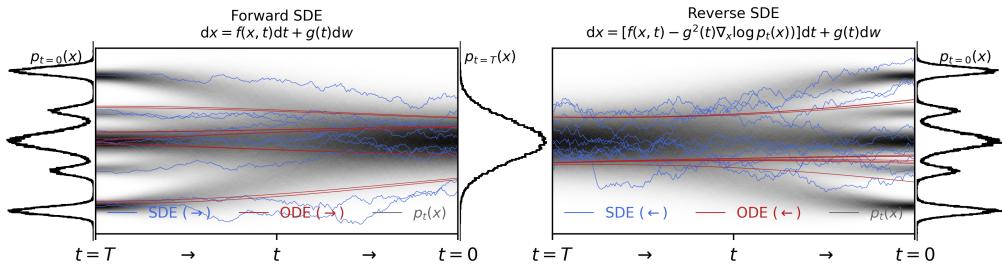


Figure 1: A diagram showing how to map data to a noise distribution (the prior) with an SDE, and reverse this SDE for generative modeling. One can also reverse the associated probability flow ODE, which yields a deterministic reverse process. Both the reverse-time SDE and probability flow ODE can be obtained by estimating the score.

46 Diffusion

47 Diffusion in the context of generative modelling describes the process of adding small amounts
48 of noise sequentially to samples of data x (Sohl-Dickstein et al., 2015). A generative model
49 for the data arises from training a neural network to reverse this process by subtracting the
50 noise added to the data.

51 Score-based diffusion models (Song, Sohl-Dickstein, et al., 2021) model a forward diffusion
52 process with Stochastic Differential Equations (SDEs) of the form

$$dx_t = f(x_t, t)dt + g(t)d\omega_t,$$

53 where $f(x_t, t)$ is a vector-valued function called the drift coefficient, $g(t)$ is the diffusion
54 coefficient and $d\omega_t$ is a sample of noise $d\omega_t \sim \mathcal{G}[d\omega_t | \mathbf{0}, \mathbf{I}_{x_t}]$. This equation describes the
55 infinitely many samples of noise along the diffusion time t that perturb the data. The diffusion
56 path, defined by the SDE, begins at $t = 0$ and ends at $T = 0$ where the resulting distribution
57 is then a multivariate Gaussian with mean zero and covariance \mathbf{I} . The code implements various
58 SDEs known in the diffusion model literature.

59 The reverse of the SDE, mapping from multivariate Gaussian samples $x(T)$ to samples of data
60 x_0 , is of the form

$$dx_t = [f(x_t, t) - g^2(t)\nabla_{x_t} \log p_t(x_t)]dt + g(t)d\omega_t,$$

61 where the score function $\nabla_{x_t} \log p_t(x_t)$ is substituted with a neural network $s_\theta(x(t), t)$ for
62 the sampling process. The network is fit by score-matching (Hyvärinen, 2005; Vincent, 2011)
63 across the time span $[0, T]$. This network predicts the noise added to the image at time
64 t with the forward diffusion process, in accordance with the SDE, and removes it. With a
65 data-dimensional sample of Gaussian noise from the prior $p_T(x)$ (see Figure 1) one can reverse
66 the diffusion process to generate data.

67 The reverse SDE may be solved with Euler-Murayama sampling (Song, Sohl-Dickstein, et al.,
68 2021) (or other annealed Langevin sampling methods) which is featured in the code.

69 Likelihood calculations with diffusion models

70 Many of the applications of generative models depend on being able to calculate the likelihood
 71 of data. Song, Sohl-Dickstein, et al. (2021) show that any SDE may be converted into an
 72 ordinary differential equation (ODE) without changing the distributions $p_t(\mathbf{x}_t)$, defined by the
 73 SDE, from which the noise is sampled from in the diffusion process (denoted $p_t(x)$ and shown
 74 in grey in Figure 1). This ODE is known as the probability flow ODE (Song, Sohl-Dickstein,
 75 et al., 2021; Song, Durkan, et al., 2021) and is written

$$d\mathbf{x}_t = [f(\mathbf{x}_t, t) - g^2(t)\nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t)]dt = f'(\mathbf{x}_t, t)dt.$$

76 This ODE can be solved with an initial-value problem to sample new data or estimate its density.
 77 Starting with a data point $\mathbf{x}_0 \sim p(\mathbf{x}) = p_0(\mathbf{x}_0)$, this point is mapped along the probability
 78 flow ODE path (see the right-hand side of Figure 1) to a sample from the multivariate Gaussian
 79 prior $\mathbf{x}_T \sim p_T(\mathbf{x}_T)$. This inherits the formalism of continuous normalising flows (Chen et al.,
 80 2019; Grathwohl et al., 2018) without the expensive ODE simulations used to train these
 81 models - allowing for a likelihood estimate based on diffusion models (Song, Durkan, et al.,
 82 2021). The initial value problem provides a solution \mathbf{x}_T and the change in probability along the
 83 path $\Delta = \log p_0(\mathbf{x}_0) - \log p_T(\mathbf{x}_T)$ where $p_T(\mathbf{x}_T)$ is a simple multivariate Gaussian distribution.
 84 Various ODE solvers of different orders are available (for a user to balance speed and accuracy
 85 of sampling) which are provided by diffrax (Kidger, 2022).

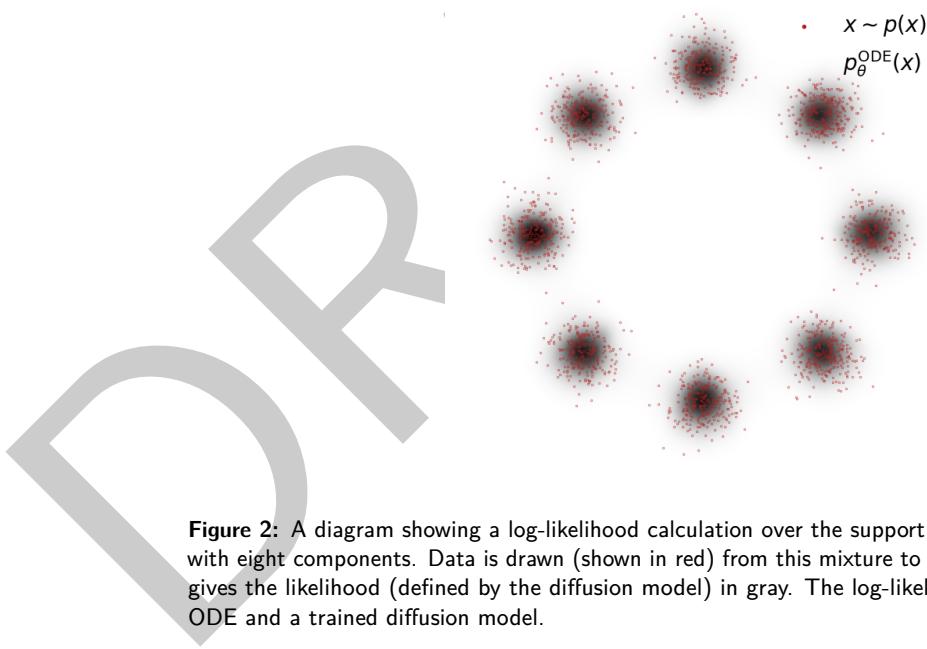


Figure 2: A diagram showing a log-likelihood calculation over the support of a Gaussian mixture model with eight components. Data is drawn (shown in red) from this mixture to train the diffusion model that gives the likelihood (defined by the diffusion model) in gray. The log-likelihood is calculated using the ODE and a trained diffusion model.

86 The likelihood estimate under a score-based diffusion model is estimated by solving the
 87 change-of-variables equation for continuous normalising flows.

$$\frac{\partial}{\partial t} \log p_t(\mathbf{x}_t) = \nabla_{\mathbf{x}_t} \cdot f(\mathbf{x}_t, t),$$

88 which gives the log-likelihood of a single datapoint \mathbf{x}_0 as

$$\log p(\mathbf{x}_0) = \log p(\mathbf{x}_T) + \int_{t=0}^{t=T} dt \nabla_{\mathbf{x}_t} \cdot f(\mathbf{x}_t, t).$$

89 The code implements these calculations also for the Hutchinson trace estimation method and
 90 (1990) that reduces the computational expense of the estimate. Figure 2 shows an example of
 91 a data-likelihood calculation using a trained diffusion model with the ODE associated from an
 92 SDE.

93 **Implementations and future work**

94 Diffusion models are defined in sbgm via a score-network model s_θ and an SDE. All the
 95 available SDEs (variance exploding (VE), variance preserving (VP) and sub-variance preserving
 96 (SubVP) (Song, Sohl-Dickstein, et al., 2021)) in the literature of score-based diffusion models
 97 are available. We provide implementations for UNet (Ronneberger et al., 2015), Diffusion
 98 Transformers (Peebles & Xie, 2023b), MLP-Mixer (Tolstikhin et al., 2021) and Residual
 99 Network (He et al., 2015) models which are state-of-the-art for diffusion tasks. It is possible to
 100 fit score-based diffusion models to a conditional distribution $p(\mathbf{x}|\pi, \mathbf{y})$ where in typical inverse
 101 problems \mathbf{y} would be an image and π a set of parameters in a physical model for the data
 102 (Batzolis et al., 2021) (e.g. to solve inverse problems). The code is compatible with any model
 103 written in the equinox (Kidger & Garcia, 2021) framework. We recently extended the code
 104 to provide transformer-based diffusion models (Peebles & Xie, 2023a) and plan to extend to
 105 latent diffusion models (Rombach et al., 2022) and flow matching (Lipman et al., 2023).

106 **Acknowledgements**

107 We thank the developers of the packages jax (Bradbury et al., 2018), optax (DeepMind et
 108 al., 2020), equinox (Kidger & Garcia, 2021) and diffraex (Kidger, 2022) for their work and
 109 for making their code available to the community.

110 **References**

- 111 and, M. F. H. (1990). A stochastic estimator of the trace of the influence matrix for laplacian
 112 smoothing splines. *Communications in Statistics - Simulation and Computation*, 19(2),
 113 433–450. <https://doi.org/10.1080/03610919008812866>
- 114 Andrews, A., Jasche, J., Lavaux, G., & Schmidt, F. (2023). Bayesian field-level inference of
 115 primordial non-gaussianity using next-generation galaxy surveys. *Monthly Notices of the
 116 Royal Astronomical Society*, 520(4), 5746–5763. <https://doi.org/10.1093/mnras/stad432>
- 117 Batzolis, G., Stanczuk, J., Schönlieb, C.-B., & Etmann, C. (2021). *Conditional image
 118 generation with score-based diffusion models*. <https://arxiv.org/abs/2111.13606>
- 119 Bradbury, J., Frostig, R., Hawkins, P., Johnson, M. J., Leary, C., Maclaurin, D., Necula, G.,
 120 Paszke, A., VanderPlas, J., Wanderman-Milne, S., & Zhang, Q. (2018). *JAX: Composable
 121 transformations of Python+NumPy programs* (Version 0.3.13). <http://github.com/jax-ml/jax>
- 122 Cabezas, A., Corenflos, A., Lao, J., & Louf, R. (2024). *BlackJAX: Composable Bayesian
 123 inference in JAX*. <https://arxiv.org/abs/2402.10797>
- 124 Chen, R. T. Q., Rubanova, Y., Bettencourt, J., & Duvenaud, D. (2019). *Neural ordinary
 125 differential equations*. <https://arxiv.org/abs/1806.07366>
- 126 Chung, H., Kim, J., Kim, S., & Ye, J. C. (2022). *Parallel diffusion models of operator and
 127 image for blind inverse problems*. <https://arxiv.org/abs/2211.10656>
- 128 Cranmer, K., Brehmer, J., & Louppe, G. (2020). The frontier of simulation-based inference.
 129 *Proceedings of the National Academy of Sciences*, 117(48), 30055–30062. <https://doi.org/10.1073/pnas.1912789117>

- 132 Daras, G., Dimakis, A. G., & Daskalakis, C. (2024). *Consistent diffusion meets tweedie:*
133 *Training exact ambient diffusion models with noisy data.* <https://arxiv.org/abs/2404.10177>
- 134 DeepMind, Babuschkin, I., Baumli, K., Bell, A., Bhupatiraju, S., Bruce, J., Buchlovsky, P.,
135 Budden, D., Cai, T., Clark, A., Danihelka, I., Dedieu, A., Fantacci, C., Godwin, J., Jones,
136 C., Hemsley, R., Hennigan, T., Hessel, M., Hou, S., ... Viola, F. (2020). *The DeepMind*
137 *JAX Ecosystem.* <http://github.com/google-deepmind>
- 138 Feng, B. T., & Bouman, K. L. (2024). *Variational bayesian imaging with an efficient surrogate*
139 *score-based prior.* <https://arxiv.org/abs/2309.01949>
- 140 Feng, B. T., Smith, J., Rubinstein, M., Chang, H., Bouman, K. L., & Freeman, W. T.
141 (2023). *Score-based diffusion models as principled priors for inverse imaging.* <https://arxiv.org/abs/2304.11751>
- 143 Goodfellow, I. J., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., Courville,
144 A., & Bengio, Y. (2014). *Generative adversarial networks.* <https://arxiv.org/abs/1406.2661>
- 145 Grathwohl, W., Chen, R. T. Q., Bettencourt, J., Sutskever, I., & Duvenaud, D. (2018).
146 *FFJORD: Free-form continuous dynamics for scalable reversible generative models.* <https://arxiv.org/abs/1810.01367>
- 148 He, K., Zhang, X., Ren, S., & Sun, J. (2015). *Deep residual learning for image recognition.*
149 <https://arxiv.org/abs/1512.03385>
- 150 Ho, J., Jain, A., & Abbeel, P. (2020). *Denoising diffusion probabilistic models.* <https://arxiv.org/abs/2006.11239>
- 152 Hyvärinen, A. (2005). Estimation of non-normalized statistical models by score matching.
153 *Journal of Machine Learning Research*, 6(24), 695–709. <http://jmlr.org/papers/v6/hyvarinen05a.html>
- 155 Kidger, P. (2022). *On neural differential equations.* <https://arxiv.org/abs/2202.02435>
- 156 Kidger, P., & Garcia, C. (2021). Equinox: Neural networks in JAX via callable PyTrees and
157 filtered transformations. *Differentiable Programming Workshop at Neural Information*
158 *Processing Systems 2021.*
- 159 Kingma, Diederik P., Salimans, T., Poole, B., & Ho, J. (2023). *Variational diffusion models.*
160 <https://arxiv.org/abs/2107.00630>
- 161 Kingma, Diederik P., & Welling, M. (2022). *Auto-encoding variational bayes.* <https://arxiv.org/abs/1312.6114>
- 163 Lipman, Y., Chen, R. T. Q., Ben-Hamu, H., Nickel, M., & Le, M. (2023). *Flow matching for*
164 *generative modeling.* <https://arxiv.org/abs/2210.02747>
- 165 Papamakarios, G., Nalisnick, E., Rezende, D. J., Mohamed, S., & Lakshminarayanan, B.
166 (2021). *Normalizing flows for probabilistic modeling and inference.* <https://arxiv.org/abs/1912.02762>
- 168 Peebles, W., & Xie, S. (2023a). *Scalable diffusion models with transformers.* <https://arxiv.org/abs/2212.09748>
- 170 Peebles, W., & Xie, S. (2023b). *Scalable diffusion models with transformers.* <https://arxiv.org/abs/2212.09748>
- 172 Remy, B., Lanusse, F., Jeffrey, N., Liu, J., Starck, J.-L., Osato, K., & Schrabbach, T. (2023).
173 Probabilistic mass-mapping with neural score estimation. *Astronomy & Astrophysics*,
174 672, A51. <https://doi.org/10.1051/0004-6361/202243054>
- 175 Rombach, R., Blattmann, A., Lorenz, D., Esser, P., & Ommer, B. (2022). *High-resolution*
176 *image synthesis with latent diffusion models.* <https://arxiv.org/abs/2112.10752>

- 177 Ronneberger, O., Fischer, P., & Brox, T. (2015). *U-net: Convolutional networks for biomedical*
178 *image segmentation*. <https://arxiv.org/abs/1505.04597>
- 179 Rozet, F. (2024). *Azula: Diffusion models in PyTorch* (Version 0.2.04). <https://github.com/probabilists/azula>
- 181 Sohl-Dickstein, J., Weiss, E. A., Maheswaranathan, N., & Ganguli, S. (2015). *Deep unsuper-*
182 *vised learning using nonequilibrium thermodynamics*. <https://arxiv.org/abs/1503.03585>
- 183 Song, Y., Durkan, C., Murray, I., & Ermon, S. (2021). *Maximum likelihood training of*
184 *score-based diffusion models*. <https://arxiv.org/abs/2101.09258>
- 185 Song, Y., Shen, L., Xing, L., & Ermon, S. (2022). *Solving inverse problems in medical imaging*
186 *with score-based generative models*. <https://arxiv.org/abs/2111.08005>
- 187 Song, Y., Sohl-Dickstein, J., Kingma, D. P., Kumar, A., Ermon, S., & Poole, B. (2021).
188 *Score-based generative modeling through stochastic differential equations*. <https://arxiv.org/abs/2011.13456>
- 190 Spurio Mancini, A., Piras, D., Alsing, J., Joachimi, B., & Hobson, M. P. (2022). <Scp>Cos-
191 *moPower</scp>*: Emulating cosmological power spectra for accelerated bayesian inference
192 from next-generation surveys. *Monthly Notices of the Royal Astronomical Society*, 511(2),
193 1771–1788. <https://doi.org/10.1093/mnras/stac064>
- 194 Tolstikhin, I., Houlsby, N., Kolesnikov, A., Beyer, L., Zhai, X., Unterthiner, T., Yung, J.,
195 Steiner, A., Keysers, D., Uszkoreit, J., Lucic, M., & Dosovitskiy, A. (2021). *MLP-mixer:*
196 *An all-MLP architecture for vision*. <https://arxiv.org/abs/2105.01601>
- 197 Vincent, P. (2011). A connection between score matching and denoising autoencoders. *Neural*
198 *Computation*, 23(7), 1661–1674. https://doi.org/10.1162/NECO_a_00142